



Big Data & AI Course: Fraud Detection - PySpark MLib

Presented By: Meriem MEKKI

Introduction

Introduction

This mini project is part of the Big Data and AI course. The objective of this work is to build a fraud flag classification model for loan applications using Apache Spark to handle large-scale data processing.

The project is based on a financial fraud dataset composed of two linked sources: loan applications data and customer transaction records. The workflow began with Exploratory Data Analysis (EDA). Followed by data preprocessing and feature engineering, including cleaning, encoding categorical variables, scaling numerical features, and joining multiple data sources using a common customer identifier.

Finally, several machine learning classification models were trained and evaluated in a distributed Spark environment to predict fraudulent loan applications.

Introduction

Dataset Description

This project uses a financial fraud dataset composed of two complementary CSV files:

- **loan_applications.csv**
- Contains information related to loan requests, including:
 - Applicant demographics
 - Financial and employment details
 - Loan characteristics (amount, term, purpose, etc.)
 - A binary fraud_flag indicating whether the loan application is fraudulent
- **transactions.csv**
- Provides historical transaction data for customers, including:
 - Transaction type and amount
 - Merchant and location details
 - Transaction-level fraud_flag

Both datasets share a common customer_id, which allows them to be joined to enrich loan applications with customer transactional behavior.

This structure enables more robust fraud detection by combining static applicant information with dynamic financial activity.

Exploratory Data Analysis (EDA)

Exploratory Data Analysis was conducted using Python libraries (Pandas, Matplotlib) to gain an initial understanding of the data before modeling.

The main EDA steps included:

- Dataset inspection
- Fraud label distribution
- Feature exploration
- Correlation analysis
- Outlier detection
- Data quality checks

Exploratory Data Analysis (EDA)

Dataset inspection

- Loaded and examined both loan_applications and transactions datasets
- Analyzed data types, structure, and dataset dimensions
- Verified consistency of the shared customer_id between datasets

```
: loans.dtypes
```

: application_id	object
customer_id	object
application_date	object
loan_type	object
loan_amount_requested	float64
loan_tenure_months	int64
interest_rate_offered	float64
purpose_of_loan	object
employment_status	object
monthly_income	float64
cibil_score	int64
existing_emis_monthly	float64
debt_to_income_ratio	float64
property_ownership_status	object
residential_address	object
applicant_age	int64
gender	object
number_of_dependents	int64
loan_status	object
fraud_flag	int64
fraud_type	object
dtype: object	


```
[48]: transactions.dtypes
```

[48]: transaction_id	object
customer_id	object
transaction_date	object
transaction_type	object
transaction_amount	float64
merchant_category	object
merchant_name	object
transaction_location	object
account_balance_after_transaction	float64
is_international_transaction	int64
device_used	object
ip_address	object
transaction_status	object
transaction_source_destination	object
transaction_notes	object
fraud_flag	int64
dtype: object	

Exploratory Data Analysis (EDA)

Fraud label distribution

- Analyzed the distribution of the fraud_flag in both datasets
- Visualized class imbalance using count plots
- Observed that fraudulent cases represent a minority class



Exploratory Data Analysis (EDA)

Feature exploration

- Examined unique values of categorical variables (e.g., loan type, purpose of loan, transaction type, device used, merchant)
- Explored numerical features such as transaction amount, account balance, and loan amount


```
Loan Applications DataFrame Unique Value Counts:  
application_id      50000  
customer_id         18314  
application_date    1096  
loan_type           5  
loan_amount_requested 1312  
loan_tenure_months   7  
interest_rate_offered 983  
purpose_of_loan      7  
employment_status    6  
monthly_income       1101  
cibil_score          354  
existing_emis_monthly 108  
debt_to_income_ratio 3288  
property_ownership_status 3  
residential_address  18314  
applicant_age        45  
gender               3  
number_of_dependents 5  
loan_status          4  
fraud_flag           2  
fraud_type           4  
dtype: int64
```

```
Transactions DataFrame Unique Value Counts:  
transaction_id      50000  
customer_id         18318  
transaction_date    30012  
transaction_type     10  
transaction_amount    639  
merchant_category    12  
merchant_name        35312  
transaction_location  8823  
account_balance_after_transaction 30267  
is_international_transaction 2  
device_used          4  
ip_address           50000  
transaction_status    2  
transaction_source_destination 39993  
transaction_notes     12  
fraud_flag            2  
dtype: int64
```

Exploratory Data Analysis (EDA)

Correlation analysis

- Computed correlation matrices between numerical features and the fraud label
- Visualized correlations using heatmaps to identify potentially relevant predictors



Exploratory Data Analysis (EDA)

Outlier detection

- Identified numerical columns in both datasets
- Used boxplots to detect outliers in financial variables

Data quality checks

- Checked for missing values in both datasets
- Analyzed feature cardinality using unique value counts

Preprocessing

Why preprocessing is critical

- Raw data comes from multiple sources (customer + transactions)
- Data exists at different granularities
- ML models require:
 - Numerical representations
 - Aggregated behavioral signals
 - Balanced class distributions

Objective

- Transform raw data into a clean, ML-ready dataset
- Preserve fraud-related patterns
- Ensure scalability using PySpark

Preprocessing

Tools & Frameworks

Technology stack used

- PySpark
 - Distributed processing
 - Scalable feature engineering
- Spark ML Pipeline
 - StringIndexer
 - VectorAssembler

Preprocessing

Data Loading & Integration

Combining heterogeneous data sources

- Loaded:
 - Customer / loan application data
 - Transaction-level data
- Joined datasets using:
 - `customer_id`
- Created a unified DataFrame:
 - `full_df`

Preprocessing

Why Transaction Aggregation is Needed

Problem: Transaction-level data is too granular

- Each customer can have:
 - Hundreds of transactions
- ML models expect:
 - Fixed-length feature vectors

Preprocessing

Transaction Aggregation

Behavioral features extracted

For each customer:

- `txn_count` → total number of transactions
- `total_txn_amount`
- `avg_txn_amount`
- `max_txn_amount`
- `distinct_merchant_count`

Fraud-sensitive transaction signals

- `international_txn_count`
- `failed_txn_count`

Derived features:

- `international_txn_ratio`
- `failed_txn_ratio`

Why ratios?

- Normalize behavior across customers
- Avoid bias toward high-volume users

Preprocessing

Financial Feature Engineering

Capturing financial stress & capacity

Created ratio-based features:

- Loan-to-Income Ratio
 - Measures debt pressure
- EMI Burden Ratio
 - Monthly repayment stress

Preprocessing

Age-Based Risk Engineering

Domain-driven risk assumptions

- Created binary risk flag:
 - `is_high_risk_age`
 - True if $\text{age} < 21$ or $\text{age} > 65$

Age Bucketing

Transforming age into categorical groups

Age buckets created:

- `very_young`
- `young`
- `early_career`
- `mid_career`
- `senior`

Preprocessing

Categorical Feature Handling

Identified categorical columns:

- Age group
- Employment / demographic features ...

Applied:

- StringIndexer
 - Converts categories → numeric indices
 - handleInvalid = "keep"

Preprocessing

Feature Vector Assembly

Creating the ML input format

Combined:

- Numerical features (ratios, counts, amounts)
- Indexed categorical features (*_idx)

Used:

- VectorAssembler

Output:

- Single column: features

→ Required by all Spark ML classifiers

Preprocessing

Preprocessing

Handling Class Imbalance

Why this matters in fraud detection

- Fraud cases are rare
- Random splitting can:
 - Remove fraud cases from train or test

Stratified Train/Test Split

Ensuring fair evaluation

- Used stratified sampling
- Maintained class distribution:
 - 80% training
 - 20% testing
- Both sets contain:
 - Fraud & non-fraud samples

Preprocessing

Final Preprocessing Output

What I achieved

- Customer-level dataset
- Aggregated transaction behavior
- Engineered financial & risk features
- Encoded categorical variables
- Balanced train/test datasets
- ML-ready feature vectors

Training and Evaluation

Models Trained

The following spark ML models were trained:

- Logistic Regression
- Random Forest Classifier
- Gradient-Boosted Trees (GBT)
- SVM
- Naive Bayes

All models use the same:

- Feature vector
- Train/test split
- Label column (fraud_flag)

Training and Evaluation

Training Process

- Used:
 - Stratified training dataset
- Same preprocessing pipeline applied to all models
- Each model:
 - Learns decision boundaries from customer-level features

This ensures fair comparison across models.

Training and Evaluation

Why Accuracy Is Not Enough

Problem with accuracy in fraud detection

- Fraud cases are rare
- A model can achieve:
 - 95% accuracy
 - While missing most fraud cases

Training and Evaluation

Evaluation Metrics Used

Metrics selected

- Precision
 - How many predicted frauds are truly fraud
- Recall
 - How many actual frauds are detected
- F1-score
 - Balance between precision & recall
- ROC-AUC
 - Overall discrimination ability

Training and Evaluation

Evaluation Methodology

- Models evaluated on:
 - Held-out test dataset
- Used Spark's:
 - BinaryClassificationEvaluator
- Predictions include:
 - prediction
 - probability

This ensures unbiased performance estimation.

Training and Evaluation

Key Observations

- Logistic Regression:
 - Serves as a strong baseline
- Random Forest:
 - Good balance between robustness and interpretability
- Gradient Boosted Trees (GBT)
 - Best overall performance
 - Highest fraud detection capability
 - Selected as the final model → Performed cross validation with grid search on this model